

Testing of Large Scale Distributed Software

Gerson Sunyé

Lina - Université de Nantes

<http://sunye.free.fr>

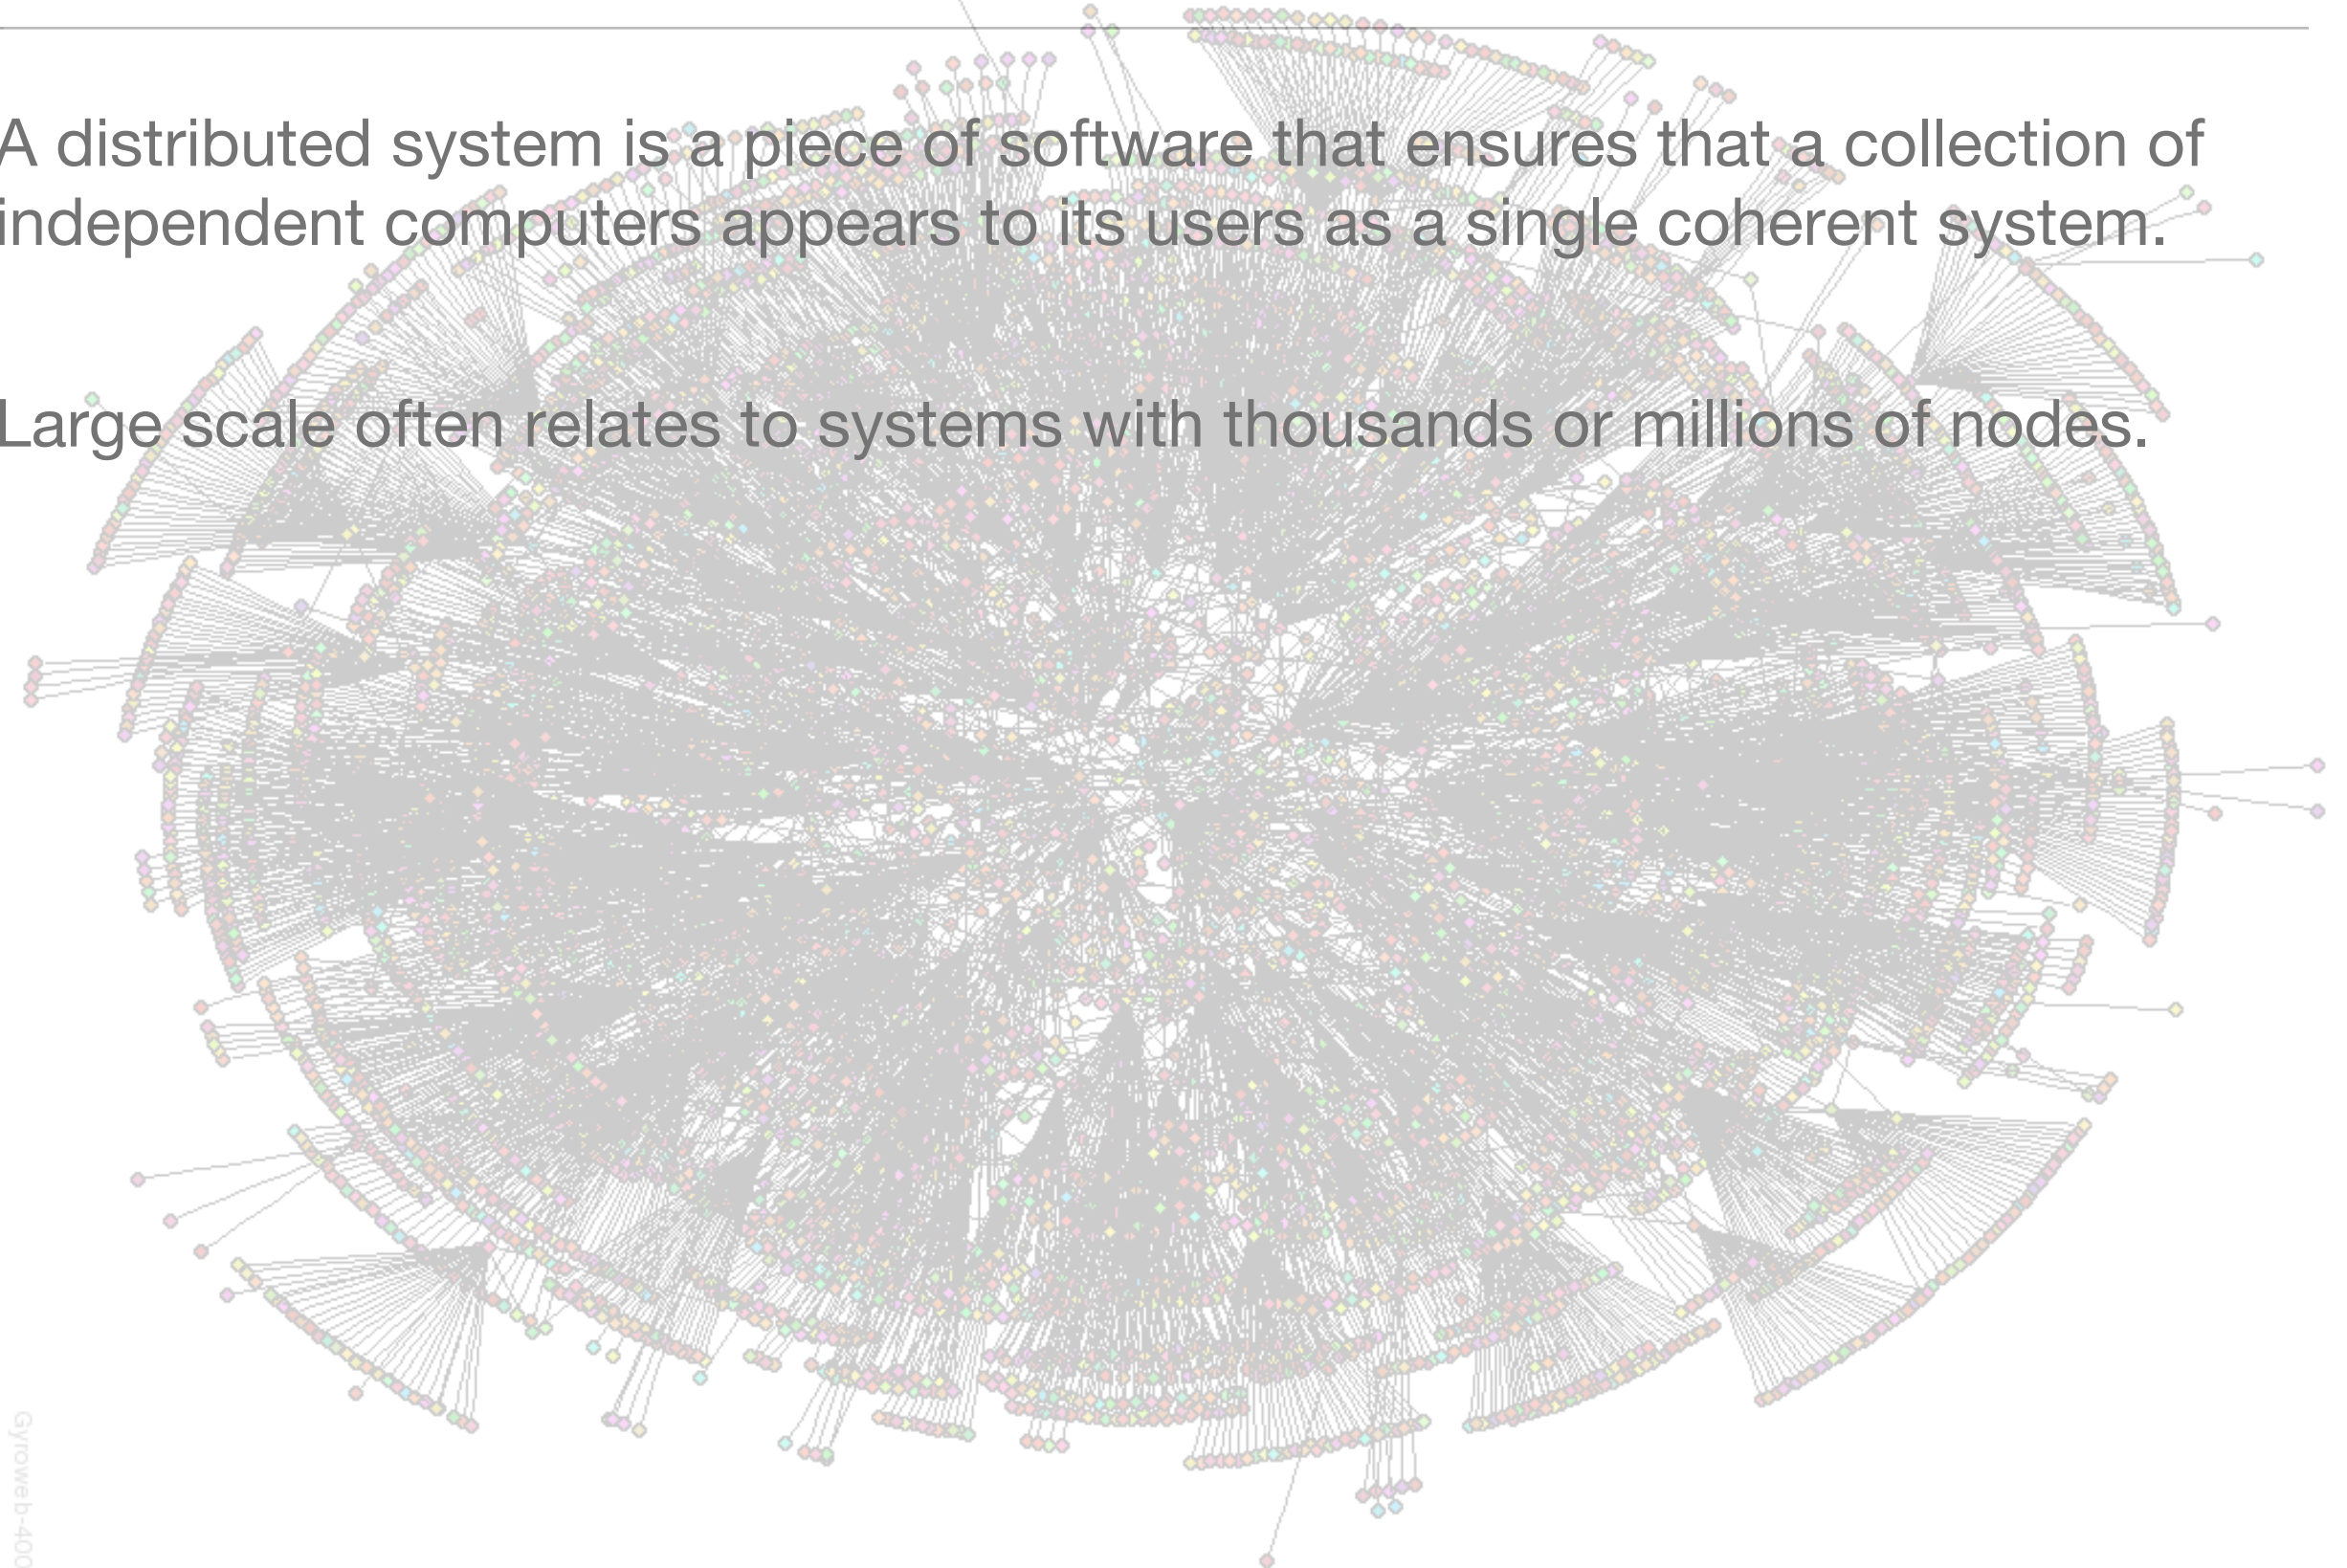
Agenda

- Large-Scale Distributed Systems.
- Characteristics of Distributed Software
- Testing of Distributed Systems
- Web Applications and Web Services.
- Peer-to-Peer Systems
- A Framework for Testing Distributed Systems
- Conclusion

Large-Scale Distributed Systems

Large-Scale Distributed (LSD) Systems

- A distributed system is a piece of software that ensures that a collection of independent computers appears to its users as a single coherent system.
- Large scale often relates to systems with thousands or millions of nodes.

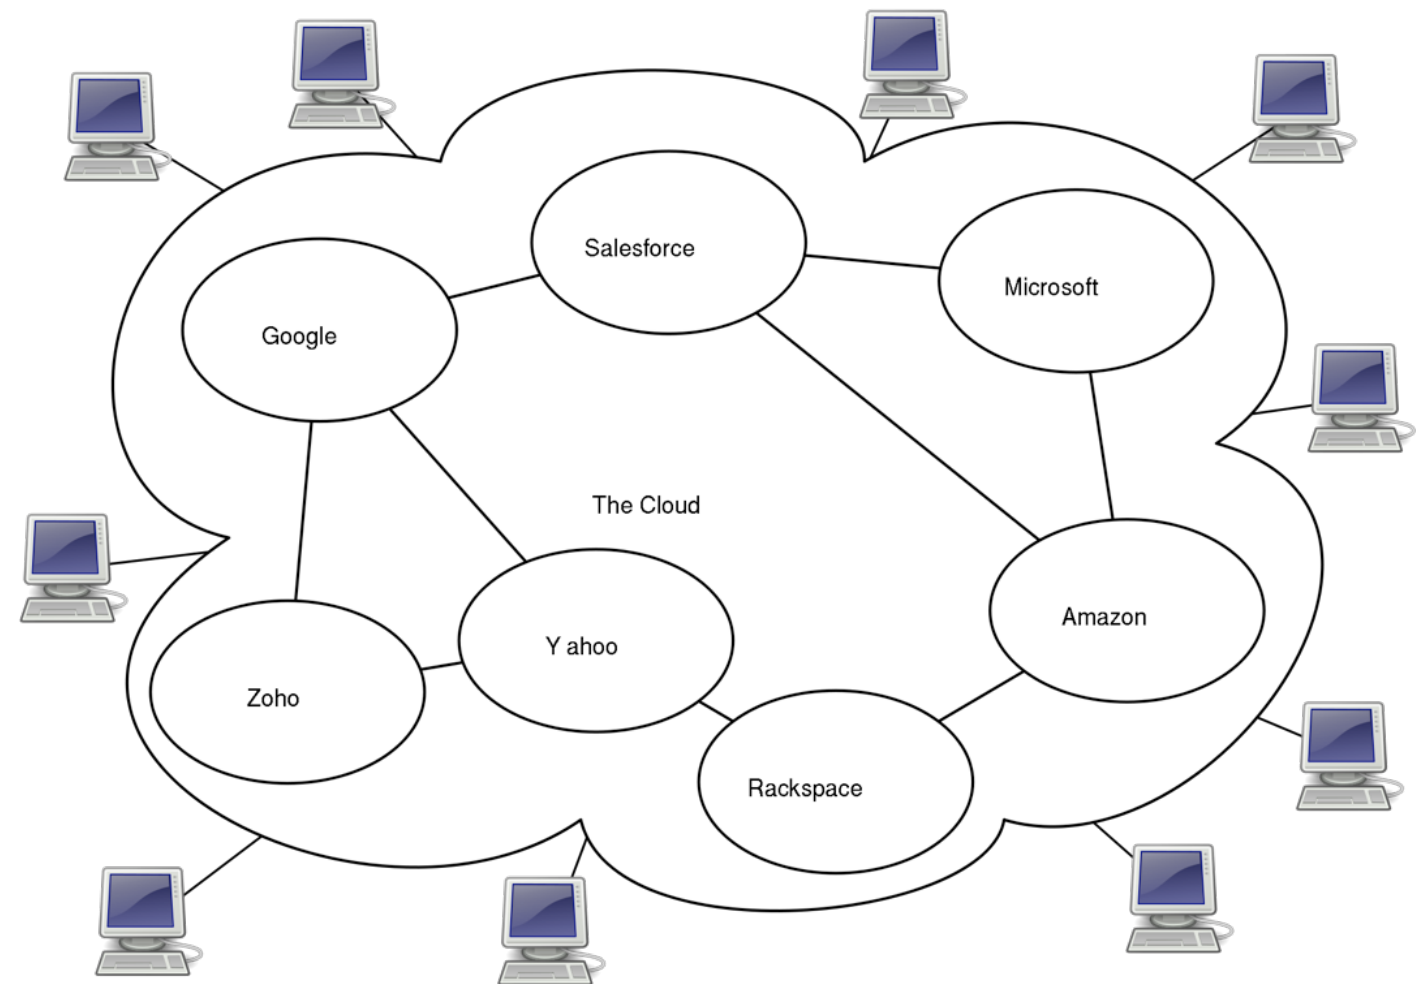


LSD Systems Examples

- Peer-to-Peer Systems: Gnutella, Napster, Skype, Joost, Distributed Hash Tables, etc.
- Cloud Computing: GoogleDocs, ThinkFree Online, etc.
- Grid Computing: Berkeley Open Infrastructure for Network Computing (BOINC), etc.
- MapReduce: Hadoop, Google, Oracle, etc.

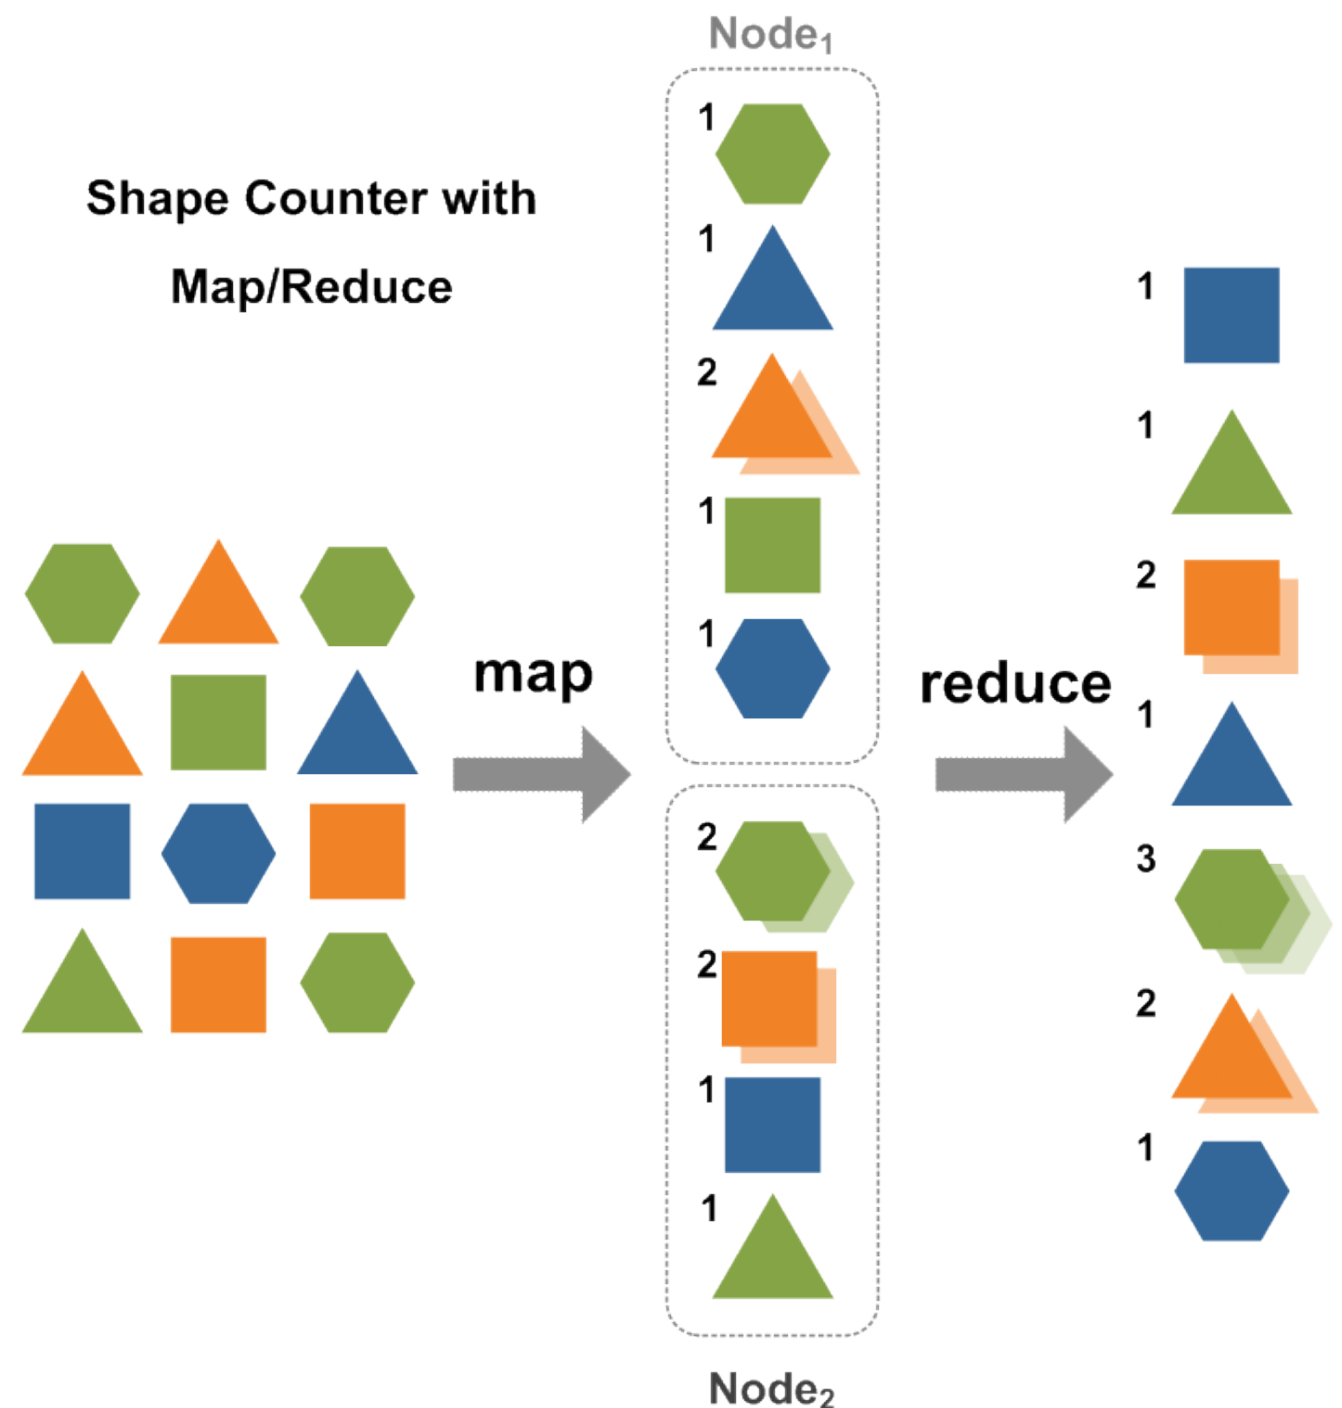
Cloud Computing

- Main concept: the computing is "in the cloud".
- Cloud service providers: Amazon, Microsoft, Google.
- Infrastructure as a service.
- Software as a service.
- Examples: Eucalyptus, OpenNebula, etc.



MapReduce

- Framework introduced by Google.
- Supports distributed computing on large data sets on clusters of computers.
- Examples: Apache Hadoop, Greenplum's, Aster Data's, etc.
- Users: Google Index, Facebook,...



Distributed Hash Table (DHT)

- Second generation P2P overlay network (Structured).
- Self-organizing, load balanced, fault tolerant.
- Scalable: guarantees the number of hops to answer a query.
- Simple interface:
 - `store(key, value)`
 - `retrieve(key) : value`

DHT Principles


- Each node is responsible for a range of keys.
- Core operation: find the node responsible for a given key.



DHT

- Node X Id association:
 - 1 node knows all other nodes (centralized).
 - All nodes know all other nodes (replication).
 - All nodes have a partial view of the whole system.

DHT Example: Chord



Why do we need DHTs ?

- Abstraction:
 - Persistence, volatility, message routing, etc.
- Applications:
 - File systems.
 - Data sharing (Gnutella index).

Characteristics of Distributed Software

How is distributed software different?

- Non-determinism;
- Third-party infrastructure;
- Partial failures;
- Time-outs;
- Dynamic nature of the structure.

Non-determinism

- Difficult to reproduce a test execution.
- The thread execution order may be affected by external programs.
- Some defects do not appear on all executions.

Third-party infrastructure

- Distributed software often rely on third-party middleware: RPC, Message Exchanging, Brokers, SOA, etc.
- Infrastructure may be self-organizing.
- Infrastructure (services) may belong to other entities.

Partial Failures

- A failure in a particular node prevents part of the system from achieving its behavior.

Time-Outs

- Time-outs are used to avoid deadlocks.
- When a response is not received within a certain time, the request is aborted.
- This is not an error.

Dynamic nature of the structure.

- Often, the architecture changes.
- Specific requests may be redirected dynamically to different nodes, depending on the load on various machines
- The number of participating nodes may vary.

Distributed Software Testing

How is Distributed Testing Different?

- Test harness deployment.
- Results retrieval.
- Execution synchronization.
- Node failure simulation.
- Number of nodes variation.

How is Distributed Testing Different?

- Level: System Testing.
- Growing need of non-functional testing: security, load, stress.

Test harness deployment

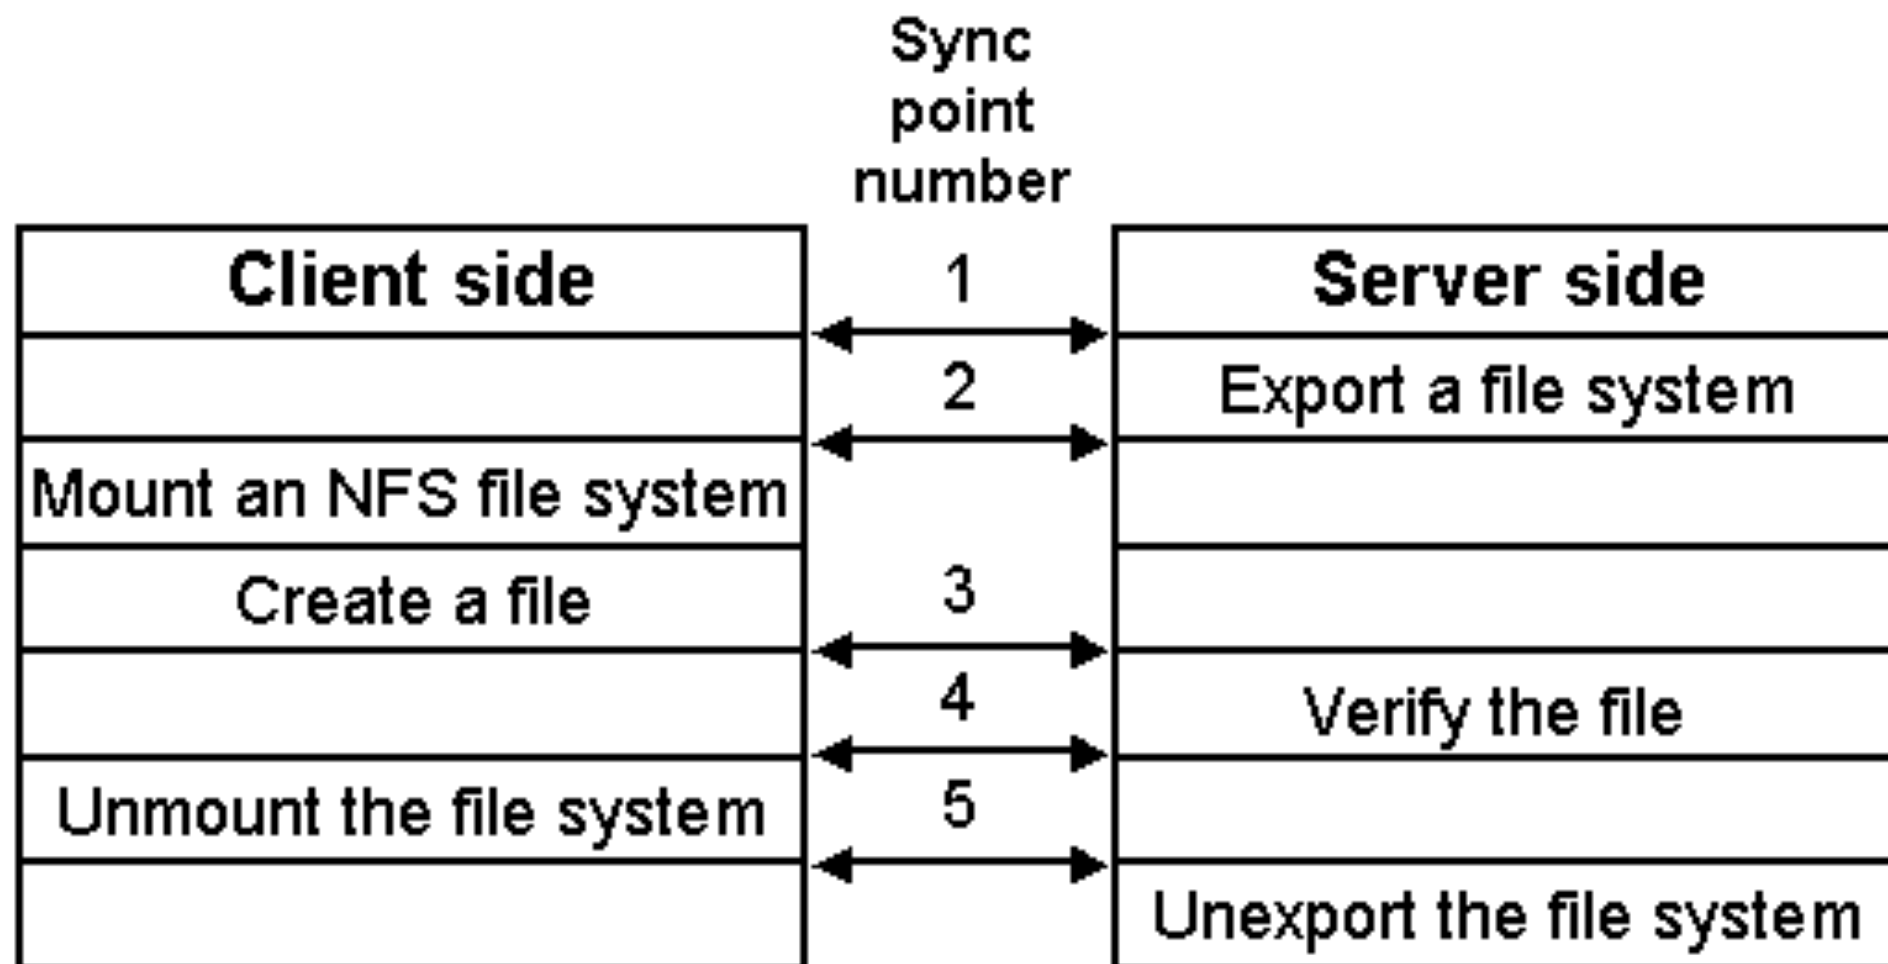
- Test harness must be deployed to every node that composes the SUT.
 - Input data.
 - Test cases.
 - Software under test.
 - Libraries.

Results retrieval

- All results must be retrieved.
- A timeline must be constructed.

Synchronization Service

- The test architecture must ensure synchronization among nodes.






Node Failure Simulation

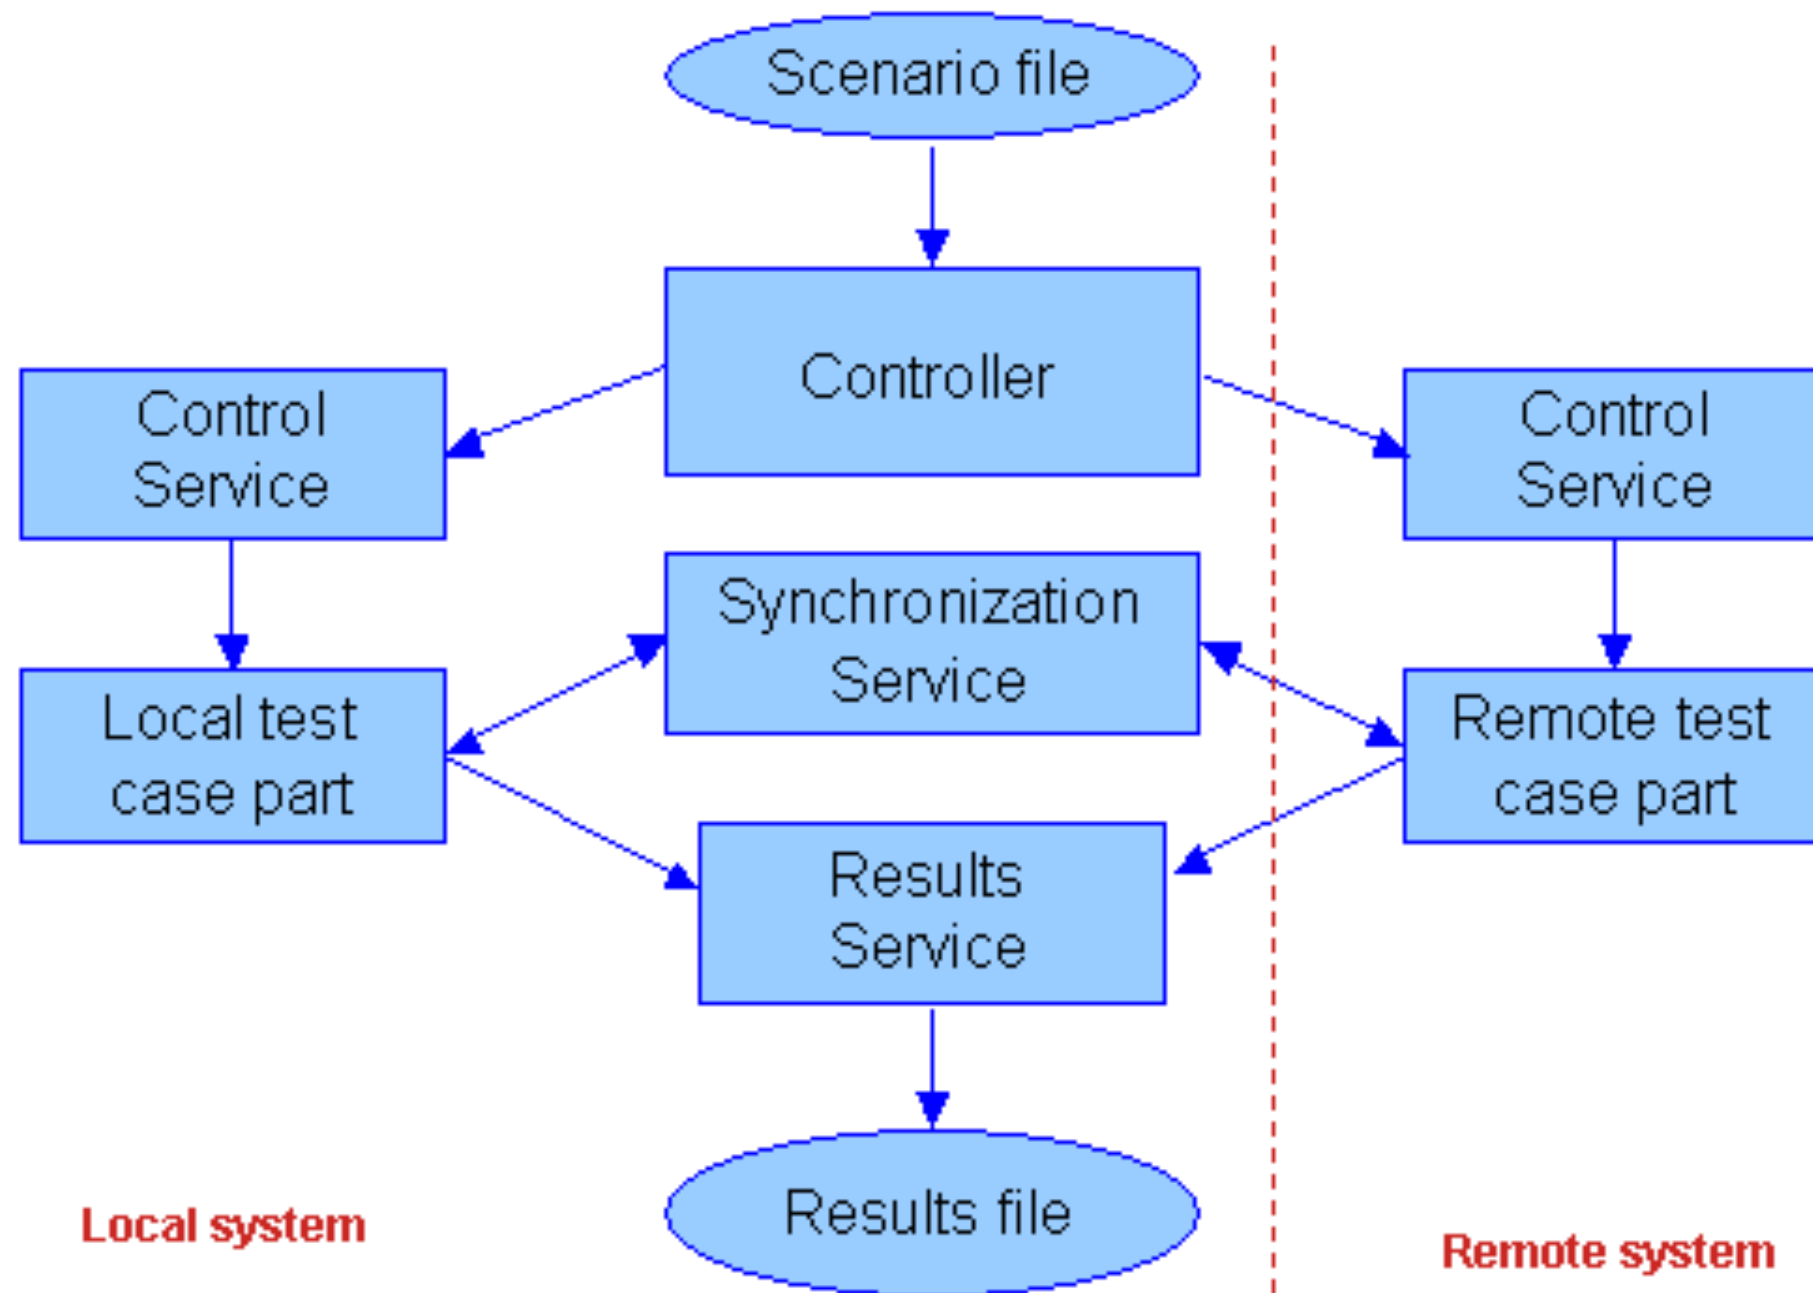
- Removing network connections.
- Shutting down a node.
- Blocking all threads of a node.
- Testing must ensure that software behaves correctly when time-outs are reached.

Number of Nodes Variation

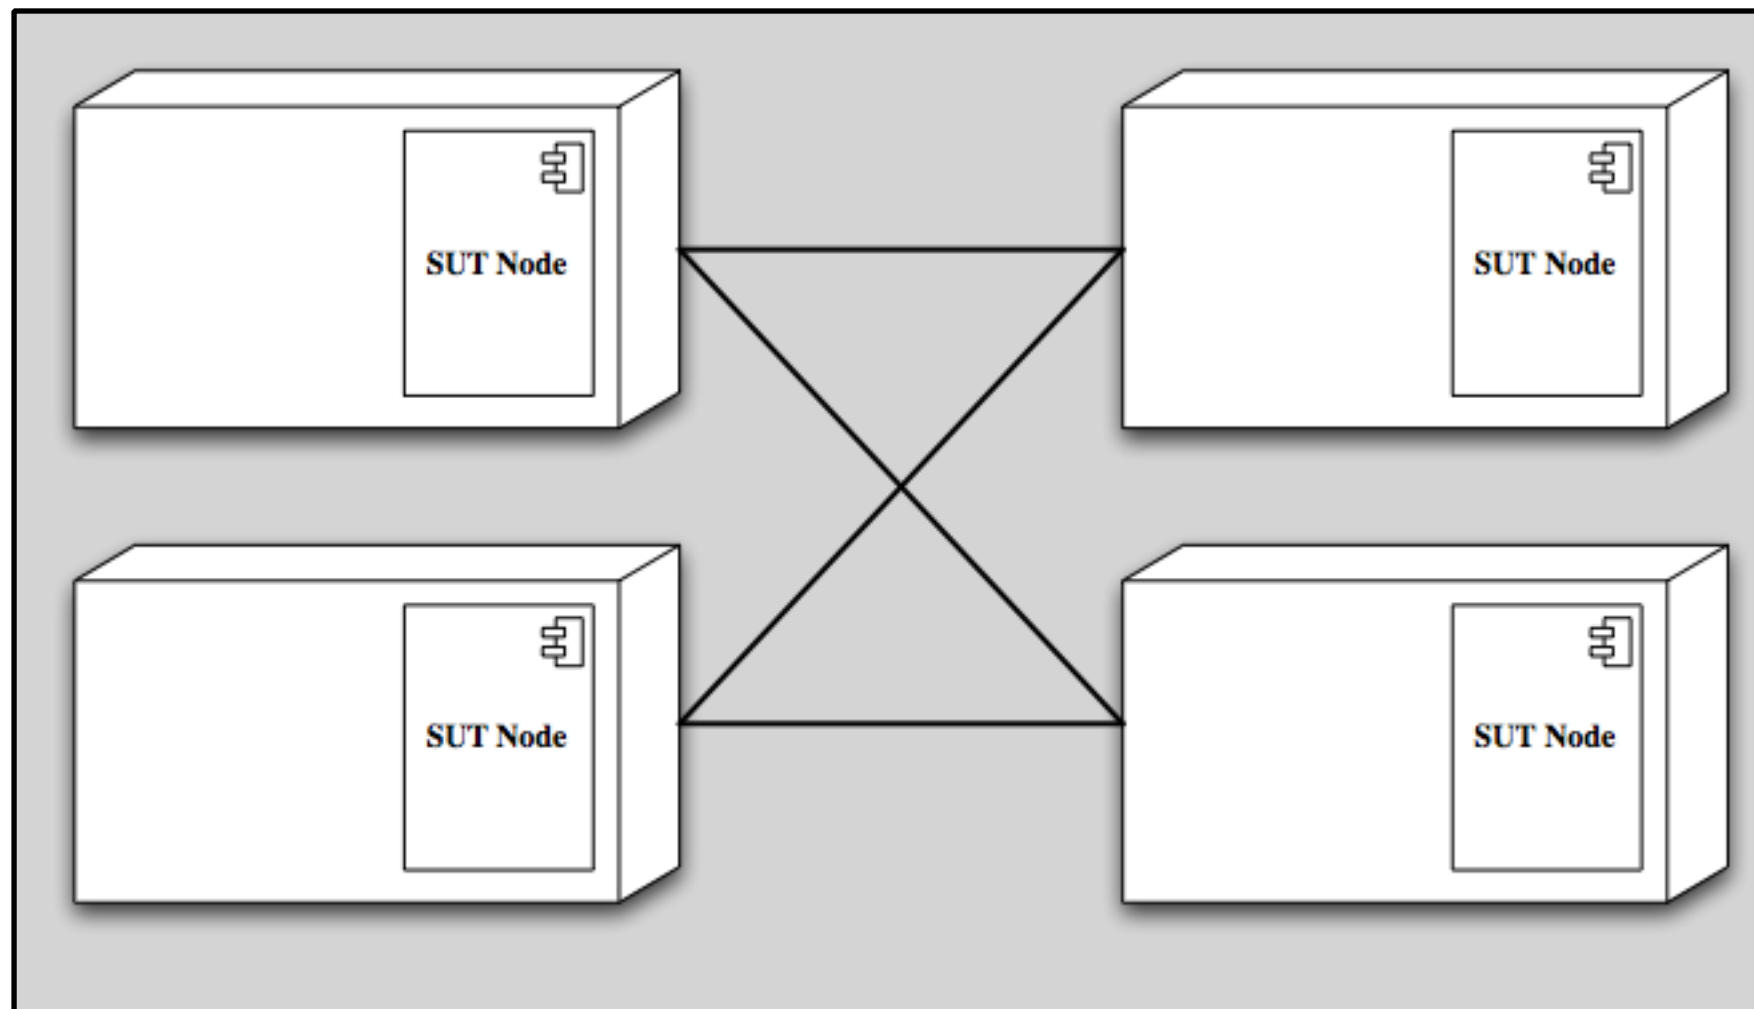
- Tests must be executed on different configurations.



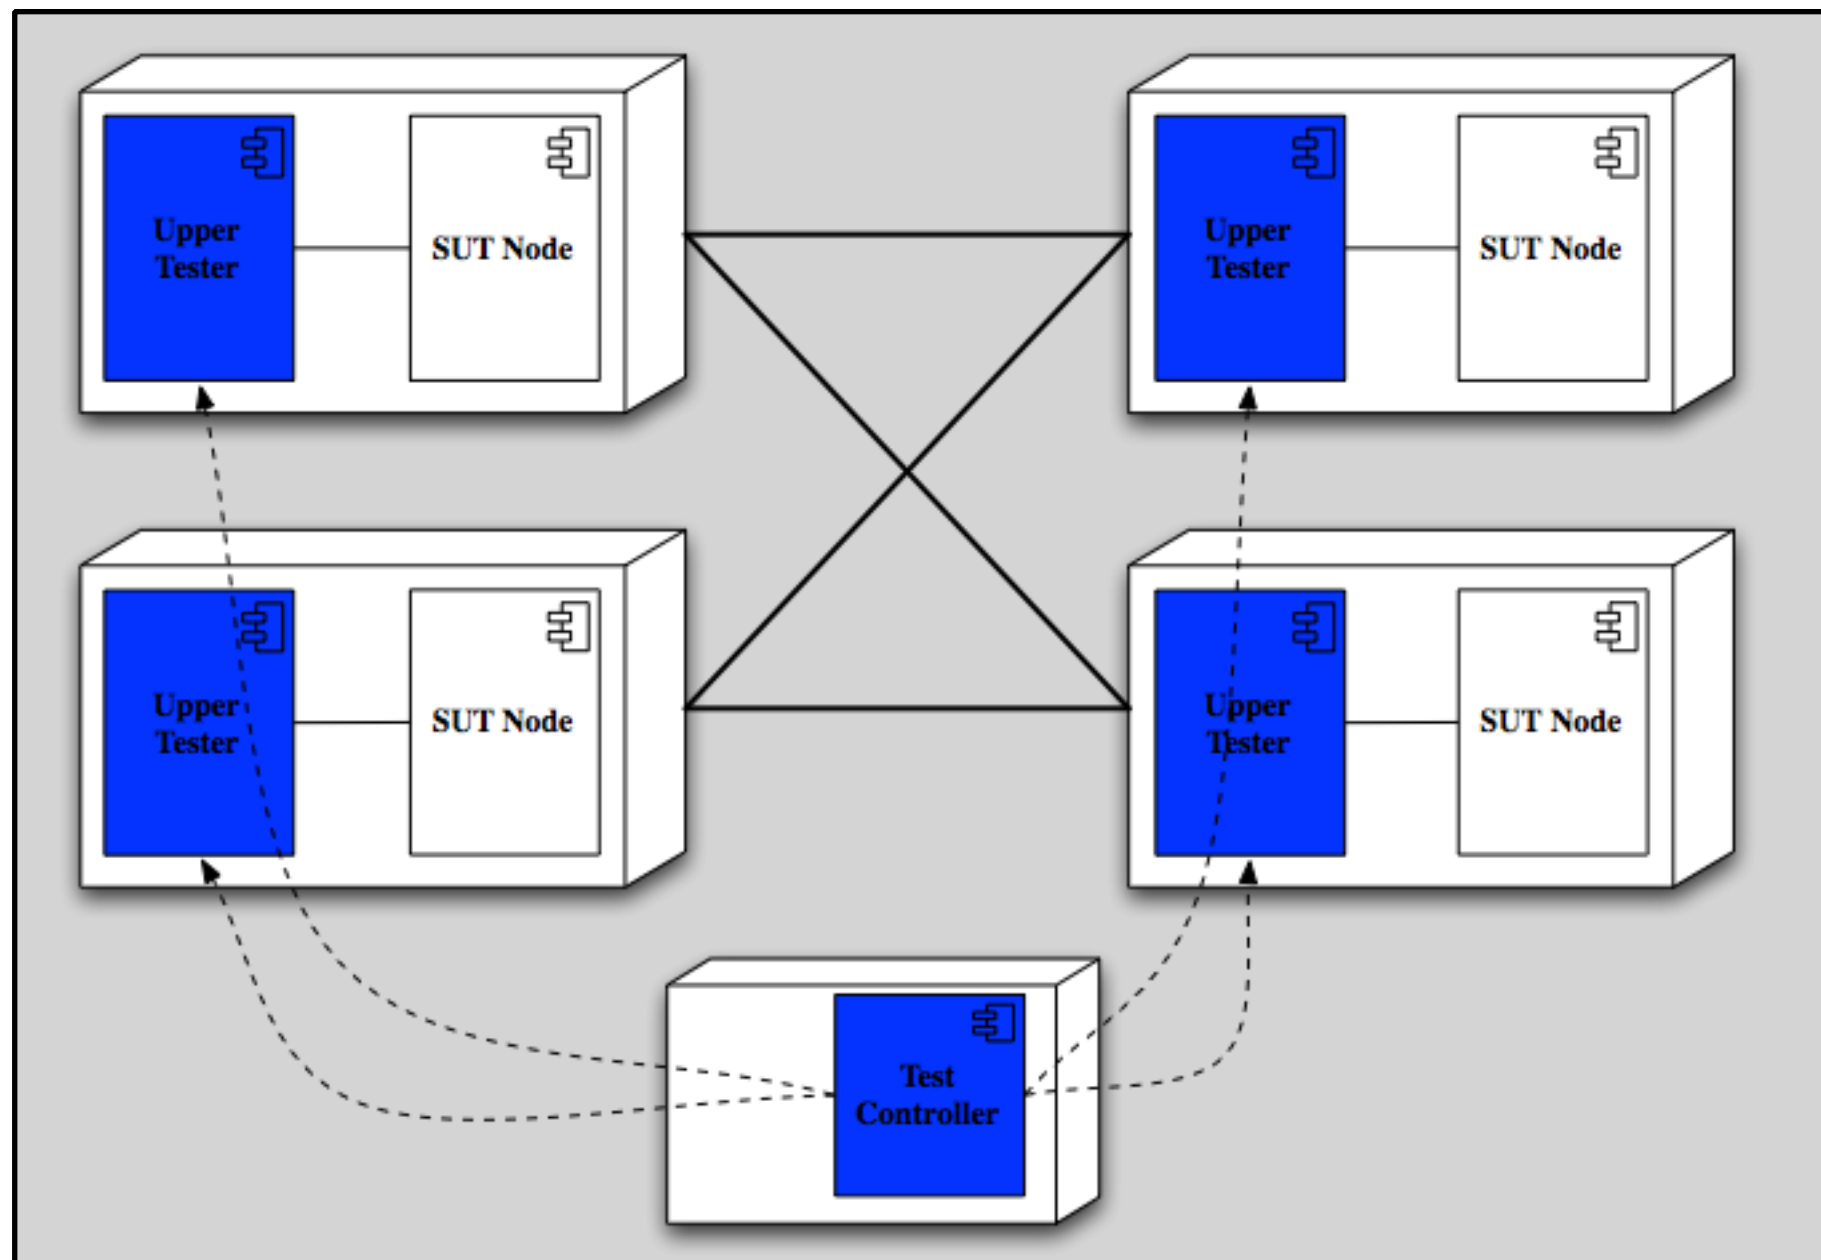
Distributed Test Architecture



A Distributed System



Distributed Test Architecture

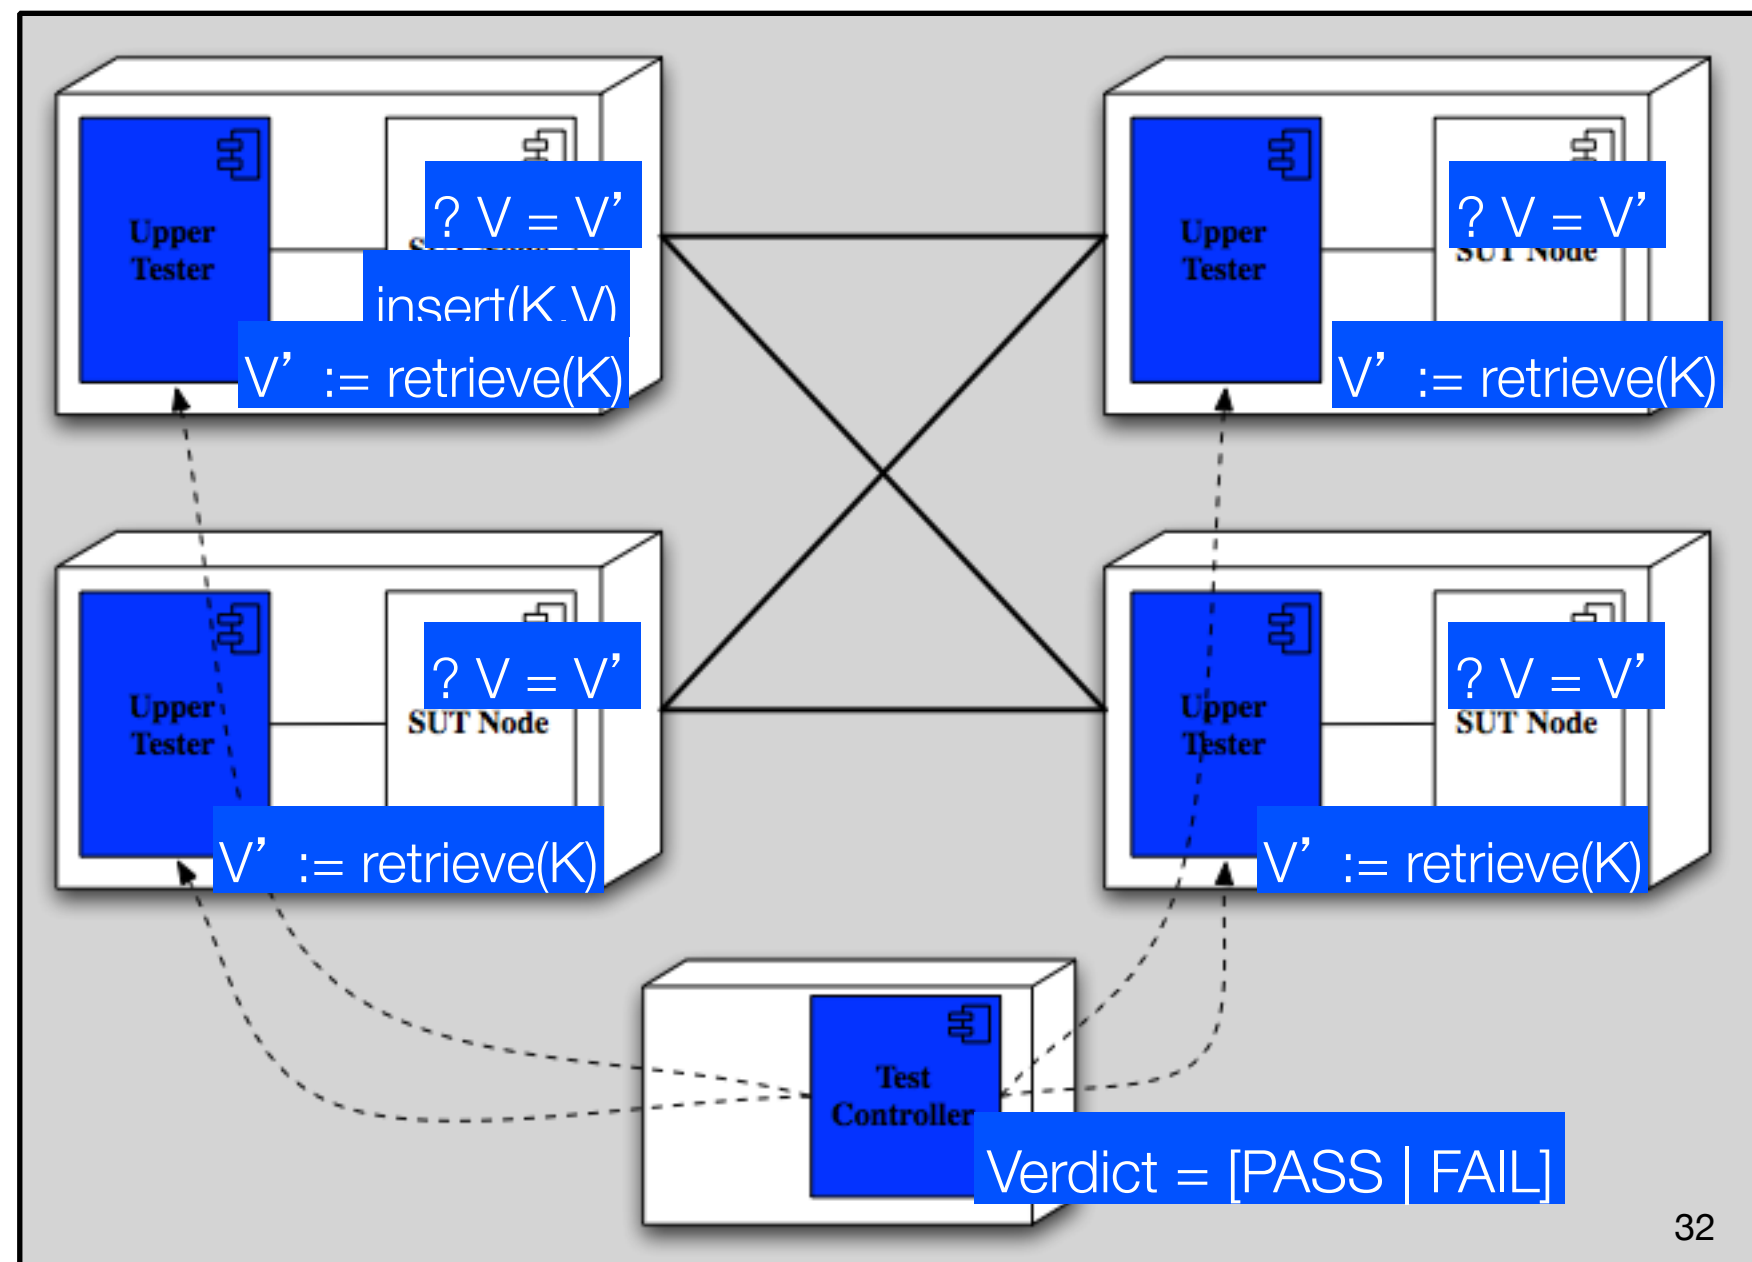


A Distributed Test Case

| Step | Action | Nodes |
|------|----------------------------|-------|
| 1 | insert(K,V) | 1 |
| 2 | $V' := \text{retrieve}(K)$ | all |
| 3 | $? V = V'$ | all |

Test Case Execution

| Step | Action | Nodes |
|------|----------------------------|-------|
| 1 | insert(K,V) | 1 |
| 2 | $V' := \text{retrieve}(K)$ | all |
| 3 | $? V = V'$ | all |



Sample Defects

- Link failure between two nodes.
- A single node failure.
- Wrong shared data access.

Web Software

Web Software

- Composed of two or more components:
 - Loosely coupled.
 - Communicate by messages.
 - Almost no shared memory.
- Concurrent and distributed.
- Technological jungle.

Web Applications

- A software deployed on the web.
- HTML User Interface.
- International (available worldwide).
- HTTP as underlying protocol.

How is Web Application Testing Different?

- Traditional graphs (control flow, call) do not apply.
- State behavior is hard to model and to describe (HTTP is stateless).
- All inputs go through the HTML UI – low controllability.
- Stress testing is crucial.

Common difficulties

- Stateless protocol: requests are independent from each other.
- Different hardware devices, programming languages, versions, etc.

New Problems

- The user interface is created dynamically.
- The flow of control may be changed:
 - back/forward, refresh
 - HTML may be rewritten!
- Input validation.
- Dynamic integration: modules can be loaded dynamically.

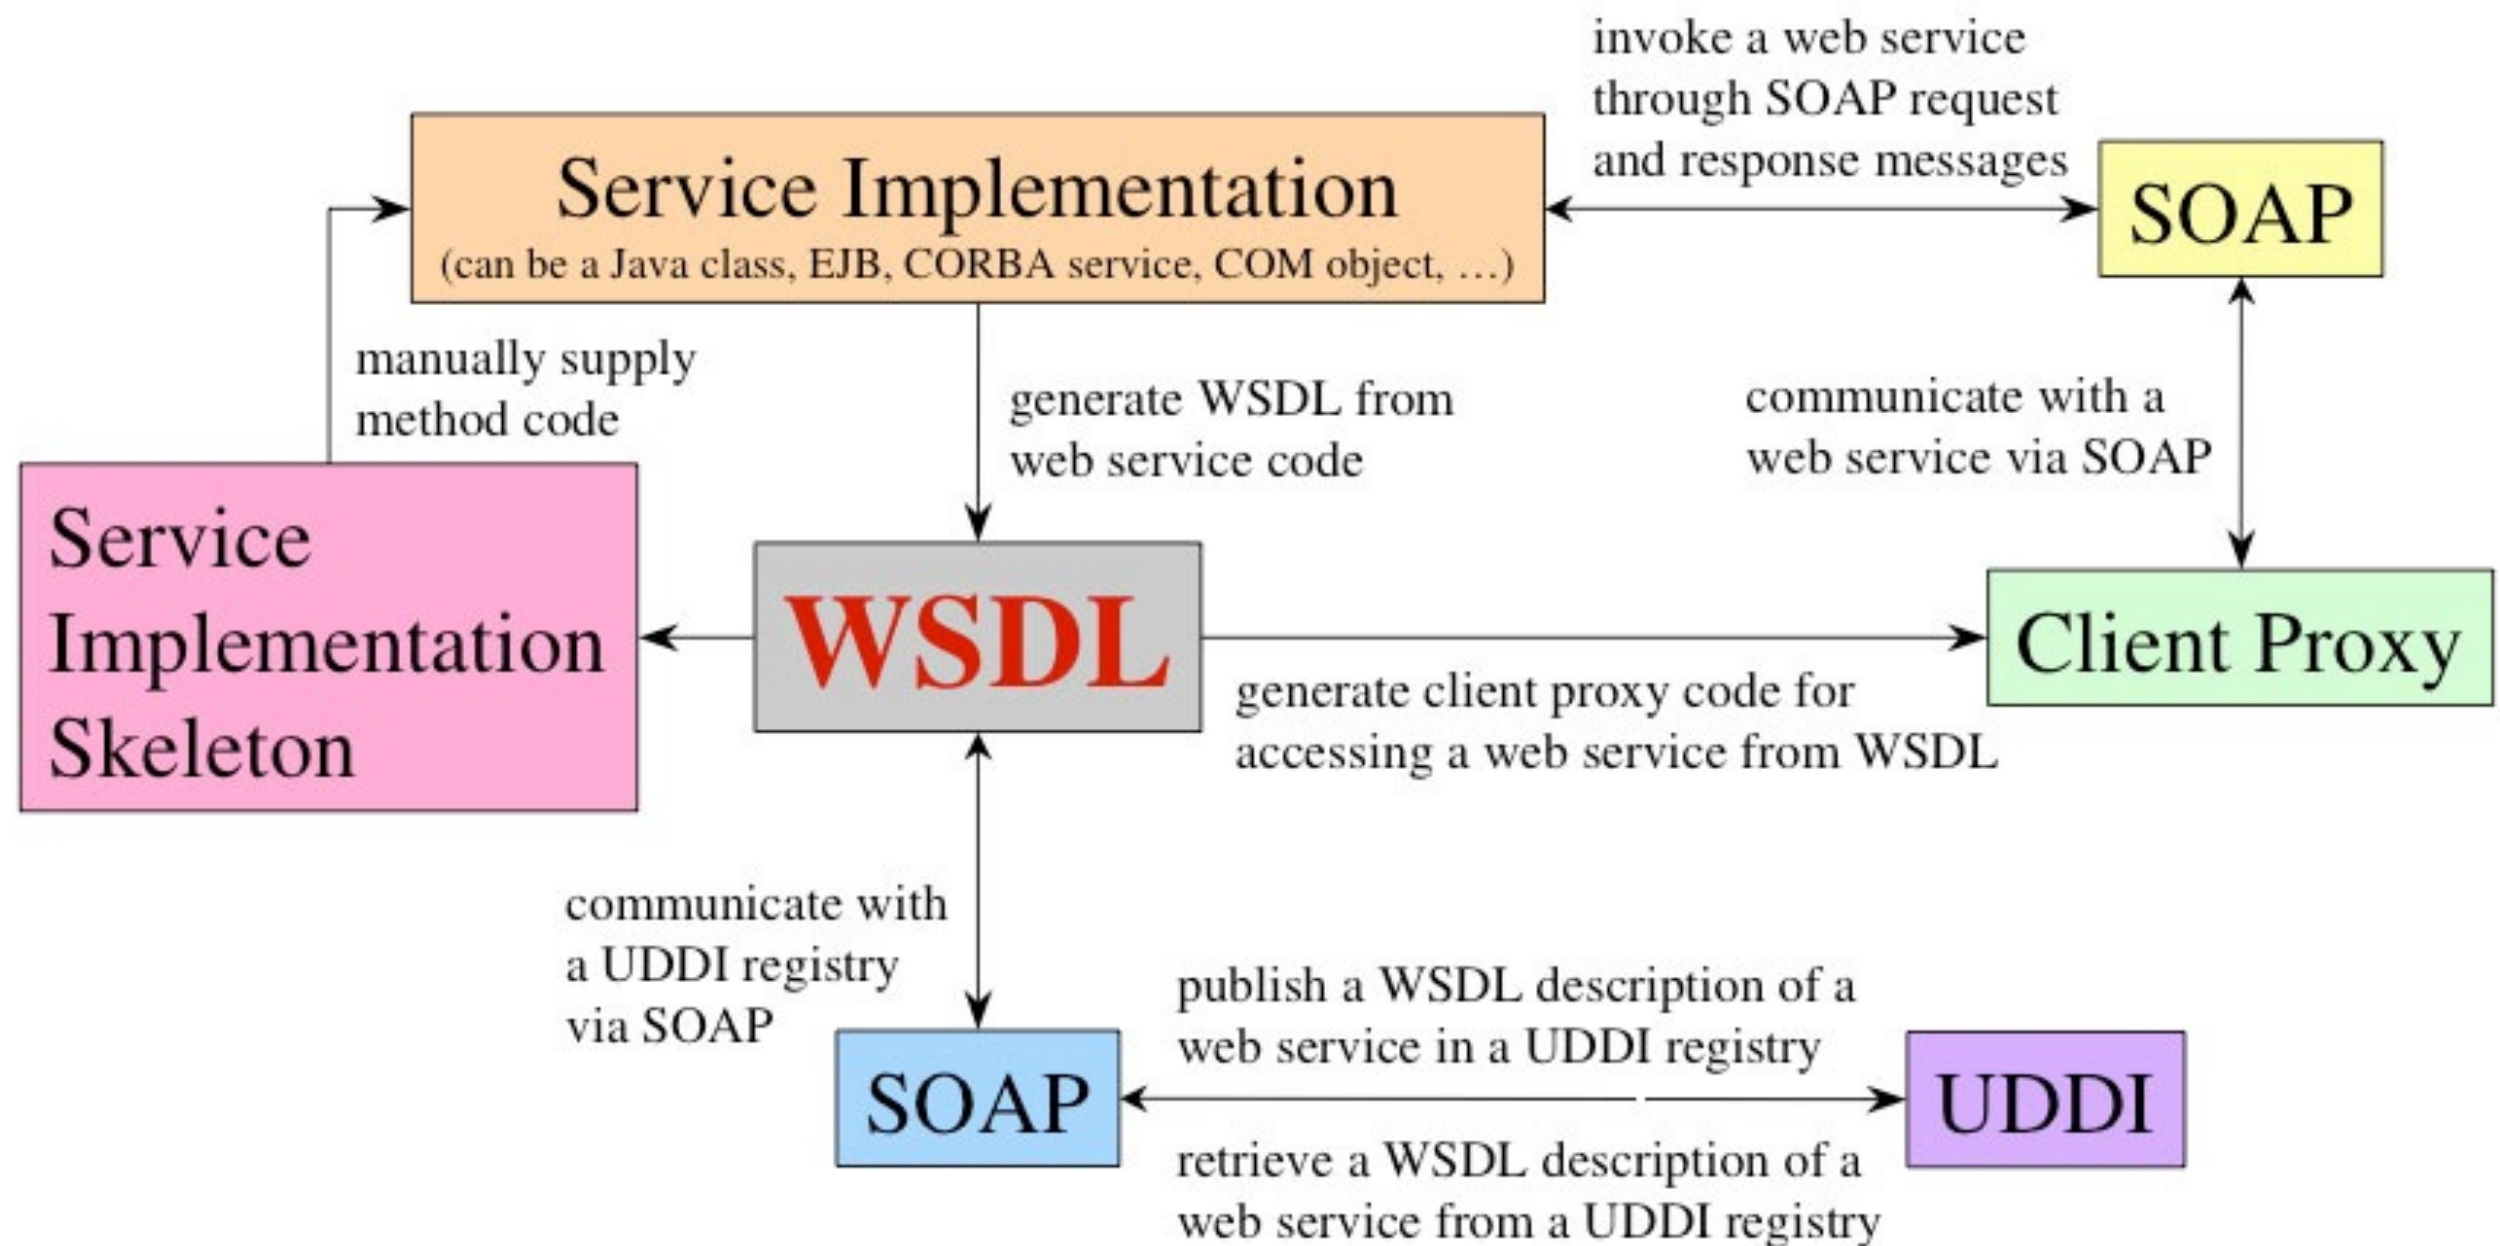
- Hard to get access to server-side state (memory, files, database) – low observability
- Not clear what logic predicates can be effectively used
- No model for mutation operators on web software

Web Services

Web Services

- A software deployed on the web that accepts remote procedure call (RPC): SOAP, XML-RPC, RESTful.
- No human interface.
- Must be published to be discovered.

Service Publication



How is Web Services Testing Different?

- Communications are often between services published by different organizations.
- The owner of a service can not control the external services it uses.
- Reconfiguration: a service can be replace by another one.

Peer-to-Peer Systems

P2P Applications

- Huge number of nodes.
- No central point of failure.
- Node volatility is a common behavior.
- Nodes are autonomous.

How is P2P Testing Different?

- Volatility must be simulated and controlled.
- Centralized services are a bottleneck:
 - Synchronization.
 - Results retrieval.

Tools

Test Tools

- SeleniumHQ
- HttpUnit
- Apache JMeter
- PeerUnit

SeleniumHQ

- Selenium IDE: Firefox plugin
- Selenium Remote Control: test runner
- Selenium Grid: RC extension for parallel testing.

HttpUnit


- Server testing.
- Emulates browser behavior.

Apache JMeter


- Performance testing.
- Not restricted to web software.
 - DBMS, LDAP, JMS, SMTP, etc.

Distributed Test Architecture

A Distributed System



Distributed Testing Architecture

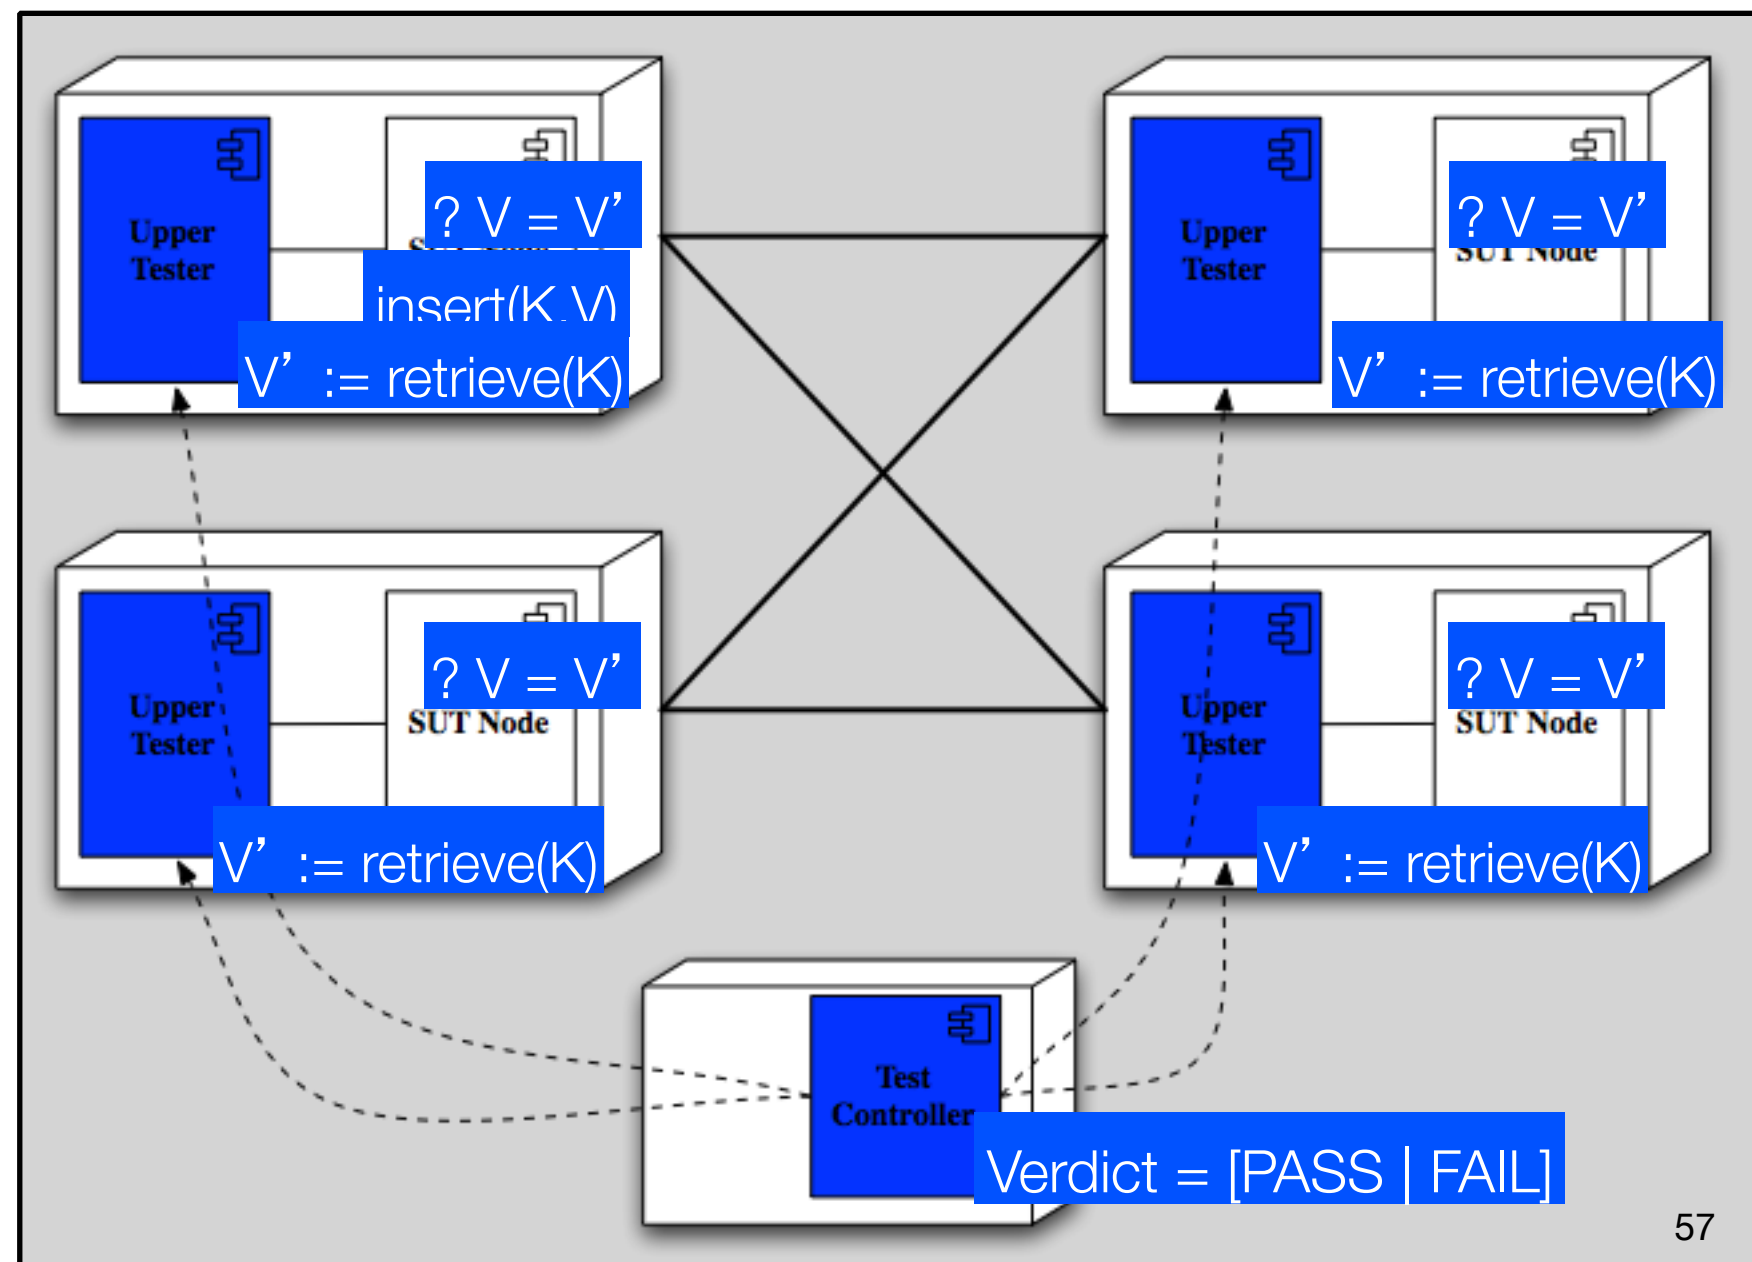


A Distributed Test Case

| Step | Action | Nodes |
|------|----------------------------|-------|
| 1 | insert(K,V) | 1 |
| 2 | $V' := \text{retrieve}(K)$ | all |
| 3 | $? V = V'$ | all |

Test Case Execution

| Step | Action | Nodes |
|------|----------------------------|-------|
| 1 | insert(K,V) | 1 |
| 2 | $V' := \text{retrieve}(K)$ | all |
| 3 | $? V = V'$ | all |



A Framework for Testing Distributed Systems

PeerUnit

- A framework for testing P2P systems;
- Written in Java 1.5;
- Open source, GPL license.

Properties

- Uses Java annotations.
- Synchronization (for $> 1,000$ testers).
- Volatility simulation.
- Shared variables.
- Distributed Architecture.

Basic Principles

| Test | PeerUnit |
|-----------|-------------|
| Test Case | Java Class |
| Test Step | Java Method |

A Simple Test Case

```
public class SimpleTest {  
    @TestStep(range = "1", timeout = 100, order = 1)  
    public void insert() {...}  
  
    @TestStep(range = "*", timeout = 100, order = 2)  
    public void retrieve() {...}  
  
    @TestStep(range = "*", timeout = 100, order = 3)  
    public void compare() {...}  
}
```

Test Execution

Step 1 - Start the Test Controller:

```
java fr.inria.peerunit.CoordinatorRunner &
```

Step 2 - Run 1 or more Testers:

```
java fr.inria.peerunit.TestRunner SimpleTest &
```

Test Results

[java] -----

[java] Test Case Verdict:

[java] Step: 1. Pass: 1. Fails: 0. Errors: 0. Inconclusive: 0. Time
elapsed: 23 msec. Average: 12 msec. Method: insert


[java] Step: 2. Pass: 7. Fails: 0. Errors: 0. Inconclusive: 0. Time
elapsed: 57 msec. Average: 11 msec. Method: retrieve

[java] Step: 3. Pass: 7. Fails: 0. Errors: 0. Inconclusive: 0. Time
elapsed: 5 msec. Average: 1 msec. Method: compare


[java] -----

PeerUnit Architecture

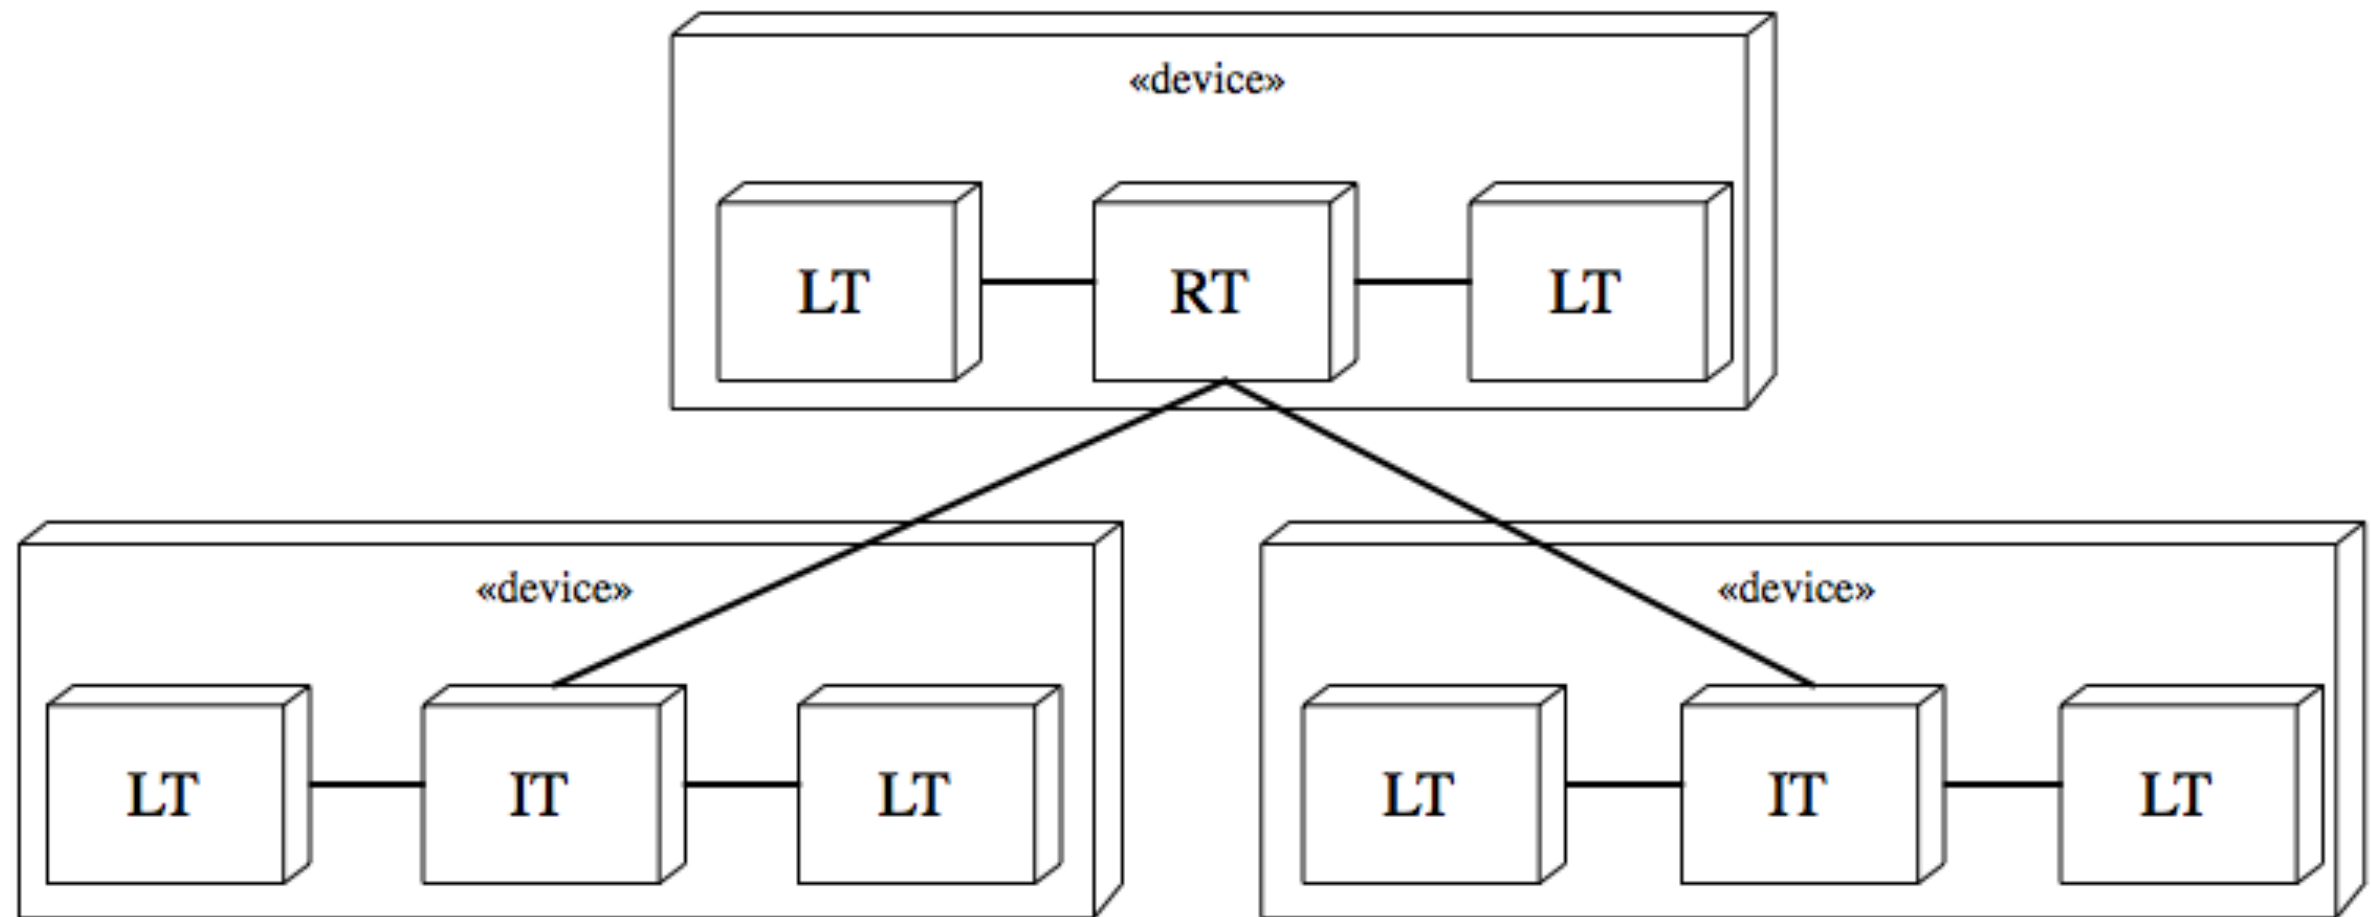
Architecture 1 : Centralized




Architecture 2: Tree



Architecture 3: Optimized Tree



Synchronization Performance



Systems already tested

- OpenChord
- FreePastry
- Kademlia
- PostgreSQL

Testing a DHT

The SUT Interface

```
public interface RemoteDHT extends Remote {  
    public void put(String k, Serializable v)  
        throws RemoteException;  
  
    public Serializable get(String k)  
        throws RemoteException;  
}
```

Test Case Summary

| Step | Action | Testers |
|------|-----------------------|---------|
| 1 | Create the (mock) DHT | 1 |
| 2 | Connect to DHT | * |
| 3 | Insert Data | 2 |
| 4 | Retrieve and Compare | * |
| 5 | Insert More Data | 1 |
| 6 | Retrieve and Compare | * |

Test Case Implementation

(RemoteDHTest.java)

Configuration

```
tester.peers=7  
tester.server=127.0.0.1  
tester.port=8181
```

```
# Log level order SEVERE,WARNING,INFO,CONFIG,FINE,FINER,FINEST  
tester.log.level=INFO  
tester.logfolder=.
```

```
# Tree Parameters  
test.treeOrder=1
```

```
# Coordination Type (0) centralized, (1) tree  
test.coordination=1
```

Test Case Execution

- Create a build file.
- Deploy the test bed.
- Execute.
- Demo.

Conclusion and References

Conclusion

- Large Scale Distributed Systems provide a new way to build software.
 - Several new technologies introducing fascinating new problems.
- Very active research area.

Open Issues

- Distributed architecture optimization.
- Lack of tools for: deployment, log analysis, graphical interface, Maven plugin, etc.
- Few users.

Bibliography

- « Testing Object-Oriented Systems: Models, Patterns, and Tools».
- Robert V. Binder.
- Addison-Wesley Object Technology Series.
- «Introduction to Software Testing».
- Paul Ammann and Jeff Offutt.
- Cambridge University Press.

« If debugging is the process of removing bugs, then programming must be the process of putting them in»

Edsger W. Dijkstra

« There' s always one more
bug»

*Lubarsky' s Law of Cybernetic
Entomology*